



IIT KHARAGPUR  
IIT GANDHINAGAR



NPTEL ONLINE  
CERTIFICATION COURSES

# Scalable Data Science

## Lecture 1: Introduction

**Sourangshu Bhattacharya**  
Computer Science and Engineering  
IIT KHARAGPUR

# In this Lecture:

- Stream processing and sketching
- Dimensionality reduction and hashing
- Frameworks for big data computation
- Scalable Machine Learning

# Stream processing and sketching

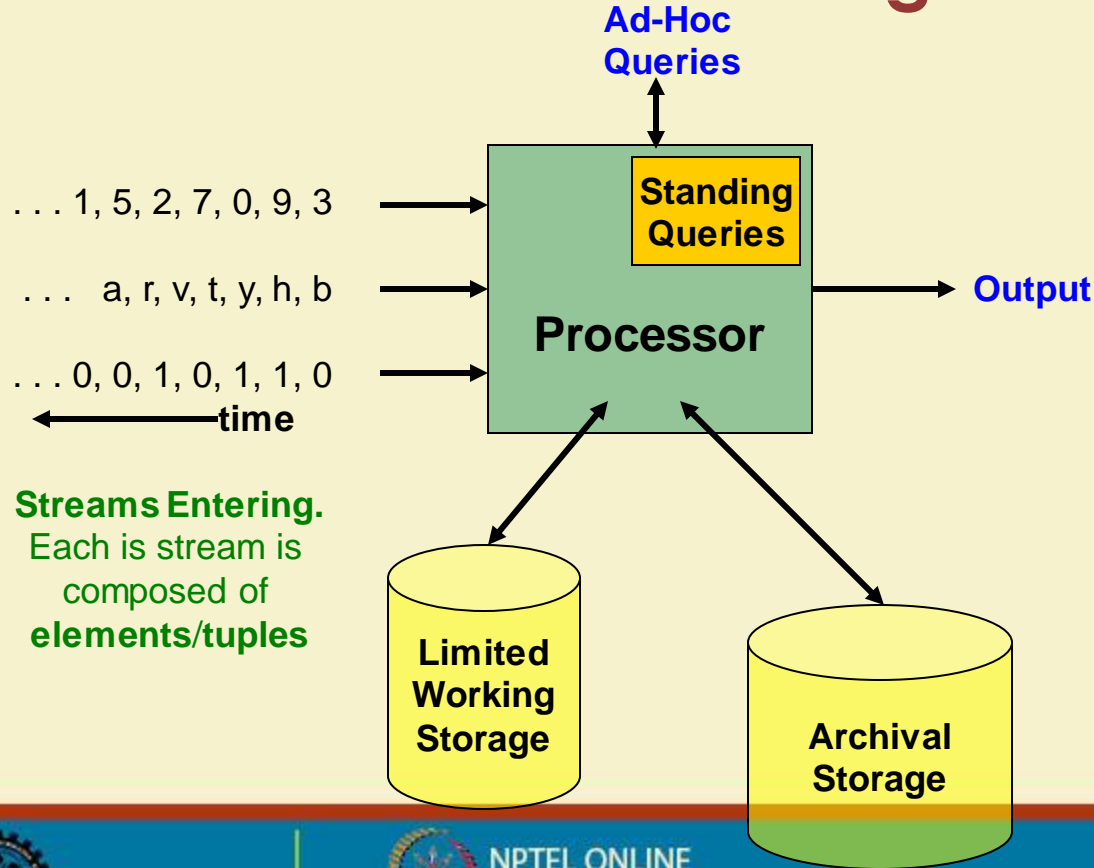
# Data Streams

- In many data mining situations, we do not know the entire data set in advance
- **Stream Management** is important when the input rate is controlled externally:
  - Google queries
  - Twitter or Facebook status updates
- We can think of the **data** as **infinite** and **non-stationary** (the distribution changes over time)

# The Stream Model

- Input **elements** enter at a rapid rate, at one or more input ports (i.e., **streams**)
  - We call **elements of the stream tuples**
- The system cannot store the entire stream accessibly
- **Q: How do you make critical calculations about the stream using a limited amount of (secondary) memory?**

# General Stream Processing Model



# Problems on Data Streams

- Types of queries one wants on answer on a data stream:
  - Sampling data from a stream
    - Construct a random sample
  - Queries over sliding windows
    - Number of items of type  $x$  in the last  $k$  elements of the stream

# Sliding Windows

- A useful model of stream processing is that queries are about a **window** of length  $N$  – the  $N$  most recent elements received
- **Amazon example:**
  - For every product  $X$  we keep 0/1 stream of whether that product was sold in the  $n$ -th transaction
  - We want to answer queries, how many times have we sold  $X$  in the last  $k$  sales




# Maintaining a fixed-size sample

- **Problem: Fixed-size sample**
- Suppose we need to maintain a random sample  $S$  of size exactly  $s$  tuples
  - E.g., main memory size constraint
- **Why?** Don't know length of stream in advance
- Suppose at time  $n$  we have seen  $n$  items
  - Each item is in the sample  $S$  with equal prob.  $s/n$

How to think about the problem: say  $s = 2$

Stream: a x c y z k c d e g...



At  $n=5$ , each of the first 5 tuples is included in the sample  $S$  with equal prob.

At  $n=7$ , each of the first 7 tuples is included in the sample  $S$  with equal prob.

**Impractical solution would be to store all the  $n$  tuples seen so far and out of them pick  $s$  at random**

# Solution: Fixed Size Sample

- **Algorithm (a.k.a. Reservoir Sampling)**
  - Store all the first  $s$  elements of the stream to  $S$
  - Suppose we have seen  $n-1$  elements, and now the  $n^{\text{th}}$  element arrives ( $n > s$ )
    - With probability  $s/n$ , keep the  $n^{\text{th}}$  element, else discard it
    - If we picked the  $n^{\text{th}}$  element, then it replaces one of the  $s$  elements in the sample  $S$ , picked uniformly at random
- **Claim:** This algorithm maintains a sample  $S$  with the desired property:
  - After  $n$  elements, the sample contains each element seen so far with probability  $s/n$

# Proof: By Induction

- We prove this by induction:

- Assume that after  $n$  elements, the sample contains each element seen so far with probability  $s/n$
- We need to show that after seeing element  $n+1$  the sample maintains the property
  - Sample contains each element seen so far with probability  $s/(n+1)$

- Base case:

- After we see  $n=s$  elements the sample  $S$  has the desired property
  - Each out of  $n=s$  elements is in the sample with probability  $s/s = 1$

# Proof: By Induction

- **Inductive hypothesis:** After  $n$  elements, the sample  $S$  contains each element seen so far with prob.  $s/n$
- **Now element  $n+1$  arrives**
- **Inductive step:** For elements already in  $S$ , probability that the algorithm keeps it in  $S$  is:

$$\left(1 - \frac{s}{n+1}\right) + \left(\frac{s}{n+1}\right)\left(\frac{s-1}{s}\right) = \frac{n}{n+1}$$

Element  $n+1$  discarded    Element  $n+1$  not discarded    Element in the sample not picked

- So, at time  $n$ , tuples in  $S$  were there with prob.  $s/n$
- Time  $n \rightarrow n+1$ , tuple stayed in  $S$  with prob.  $n/(n+1)$
- So prob. tuple is in  $S$  at time  $n+1 = \frac{s}{n} \cdot \frac{n}{n+1} = \frac{s}{n+1}$

# Problems on Data Streams

- **Types of queries one wants on answer on a data stream:**
  - **Filtering a data stream**
    - Select elements with property  $x$  from the stream
  - **Counting distinct elements**
    - Number of distinct elements in the last  $k$  elements of the stream
  - **Estimating moments**
    - Estimate avg./std. dev. of last  $k$  elements
  - **Finding frequent elements**

# Applications (1)

- **Mining query streams**
  - Google wants to know what queries are more frequent today than yesterday
- **Mining click streams**
  - A web company wants to know which of its pages are getting an unusual number of hits in the past hour
- **Mining social network news feeds**
  - E.g., look for trending topics on Twitter, Facebook


# Applications (2)

- **Sensor Networks**
  - Many sensors feeding into a central controller
- **Telephone call records**
  - Data feeds into customer bills as well as settlements between telephone companies
- **IP packets monitored at a switch**
  - Gather information for optimal routing
  - Detect denial-of-service attacks

# Dimensionality reduction



# Dimensionality Reduction



- **Assumption:** Data lies on or near a low  $d$ -dimensional subspace
- **Axes of this subspace are effective representation of the data**

# Dimensionality Reduction

- **Compress / reduce dimensionality:**
  - $10^6$  rows;  $10^3$  columns; no updates
  - Random access to any cell(s); **small error: OK**

day	We	Th	Fr	Sa	Su
customer	7/10/96	7/11/96	7/12/96	7/13/96	7/14/96
ABC Inc.	1	1	1	0	0
DEF Ltd.	2	2	2	0	0
GHI Inc.	1	1	1	0	0
KLM Co.	5	5	5	0	0
Smith	0	0	0	2	2
Johnson	0	0	0	3	3
Thompson	0	0	0	1	1

The above matrix is really “2-dimensional.” All rows can be reconstructed by scaling  $[1\ 1\ 1\ 0\ 0]$  or  $[0\ 0\ 0\ 1\ 1]$


# Why Reduce Dimensions?

## Why reduce dimensions?


- **Discover hidden correlations/topics**
  - Words that occur commonly together
- **Remove redundant and noisy features**
  - Not all words are useful
- **Interpretation and visualization**
  - Genres of movies
- **Easier storage and processing of the data**

# Locality sensitive hashing



# Scene Completion Problem



# Scene Completion Problem



# Scene Completion Problem



10 nearest neighbors from a collection of 20,000 images



# Scene Completion Problem



10 nearest neighbors from a collection of 2 million images



# A Common Metaphor

- Many problems can be expressed as finding “similar” sets:
  - Find near-neighbors in high-dimensional space
- **Examples:**
  - **Pages with similar words**
    - For duplicate detection, classification by topic
  - **Customers who purchased similar products**
    - Products with similar customer sets
  - **Images with similar features**
    - Users who visited similar websites

# Problem definition

- **Given: High dimensional data points  $x_1, x_2, \dots$** 
  - **For example:** Image is a long vector of pixel colors
$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 2 & 1 \\ 0 & 1 & 0 \end{bmatrix} \rightarrow [1 \ 2 \ 1 \ 0 \ 2 \ 1 \ 0 \ 1 \ 0]$$
- **And some distance function  $d(x_1, x_2)$** 
  - Which quantifies the “distance” between  $x_1$  and  $x_2$
- **Goal:** Find **all pairs of data points  $(x_i, x_j)$**  that are within some distance threshold  $d(x_i, x_j) \leq s$
- **Note:** Naïve solution would take  $O(N^2)$  ☹  
where  $N$  is the number of data points
- **MAGIC:** This can be done in  $O(N)$ !! How?

# Frameworks for big data computation

# MapReduce

- Much of the course will be devoted to **large scale computing for data mining**
- **Challenges:**
  - How to distribute computation?
  - Distributed/parallel programming is hard
- **Map-reduce** addresses all of the above
  - Google's computational/data manipulation model
  - Elegant way to work with big data

# Example: Language Model

- **Statistical machine translation:**
  - Need to count number of times every 5-word sequence occurs in a large corpus of documents
- **Very easy with MapReduce:**
  - **Map:**
    - Extract (5-word sequence, count) from document
  - **Reduce:**
    - Combine the counts


# Example: Host size

- **Suppose we have a large web corpus**
- Look at the metadata file
  - Lines of the form: (URL, size, date, ...)
- **For each host, find the total number of bytes**
  - That is, the sum of the page sizes for all URLs from that particular host
- **Other examples:**
  - Link analysis and graph processing
  - Machine Learning algorithms

# Scalable Machine Learning

# Big Data

- **6 Billion** web queries per day.  
~ 6 TB per day, ~ **2.5 PB** per year
- **10 Billion** display ads per day.  
~ 15 TB per day, ~ **5.5 PB** per year
- **30 Billion** text ads per day.  
~ 30 TB per day, ~ **11 PB** per year
- **150 Million** Credit card transactions per day.  
~ 150 GB per day, ~ **5.5 TB** per year
- **100 Billion** emails per day.  
~ 1 PB per day, ~ **360 PB** per year





# Machine Learning on Big Data

- **6 Billion** web queries per day.  
~ 6 TB per day, ~ **2.5 PB** per year
- **10 Billion** display ads per day.  
~ 15 TB per day, ~ **5.5 PB** per year
- **30 Billion** text ads per day.  
~ 30 TB per day, ~ **11 PB** per year
- **150 Million** Credit card transactions per day.  
~ 150 GB per day, ~ **5.5 TB** per year
- **100 Billion** emails per day.  
~ 1 PB per day, ~ **360 PB** per year
- **Ranking** search results  
Training ranking algorithms from past searches
- **Segmentation** of customers e.g. “high income male”  
View count by customer segments
- **Click through rate** estimation  
Training logistic regression
- **Fraudulent** transactions  
Anomaly detection
- **Personalised spam filtering**  
Multi-task binary classification

# Large Scale Machine Learning

- **Main question:**  
**How to efficiently train**  
(build a model/find model parameters)?
- **Auxiliary question: fast / scalable optimization**
  - Stochastic / online optimization
  - Distributed optimization.

# References:

- Jure Leskovec, Anand Rajaraman, Jeff Ullman. **Mining of Massive Datasets**. 2<sup>nd</sup> edition. - Cambridge University Press. <http://www.mmds.org/>

# Thank You!!



IIT KHARAGPUR



NPTEL ONLINE  
CERTIFICATION COURSES

Faculty Name  
Department Name