



IIT KHARAGPUR  
IIT GANDHINAGAR



NPTEL ONLINE  
CERTIFICATION COURSES

# Scalable Data Science

## Lecture 7: Bloom Filters

Anirban Dasgupta

Computer Science and Engineering

IIT GANDHINAGAR



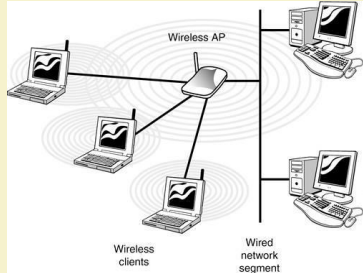


**IIT Gandhinagar**

Indian Institute of  
Technology Gandhinagar


# Querying



ISBN present in collection?



IP seen by switch?



10.0.21.102

# Solutions

- Universe  $U$ , but need to store a set of  $n$  items,  $n \ll |U|$
- Hash table of size  $m$ :
  - Space  $O(n \log |U|)$
  - Query time  $O\left(\frac{n}{m}\right)$



# Solutions

- Universe  $U$ , but need to store a set of  $n$  items,  $n \ll |U|$
- Hash table of size  $m$ :
  - Space  $O(n \log |U|)$
  - Query time  $O\left(\frac{n}{m}\right)$
- Bit array of size  $|U|$ 
  - Space =  $|U|$
  - Query time  $O(1)$




# Querying, Monte Carlo style

- In hash table construction, we used random hash functions
  - we never return incorrect answer
  - query time is a random variable
  - These are Las Vegas algorithms
- In Monte-Carlo randomized algorithms, we are allowed to return incorrect answers with (small) probability, say,  $\delta$

# Bloom filter


[Bloom, 1970]

- A bit-array  $B$ ,  $|B| = m$
- $k$  hash functions,  $h_1, h_2, \dots, h_k$ , each  $h_i \in U \rightarrow [m]$



# Bloom filter

- A bit-array  $B$ ,  $|B| = m$
- $k$  hash functions,  $h_1, h_2, \dots, h_k$ , each  $h_i \in U \rightarrow [m]$



# Operations

- *Initialize*( $B$ )
  - for  $i \in \{1, \dots, m\}$ ,  $B[i] = 0$
- *Insert* ( $B, x$ )
  - for  $i \in \{1, \dots, k\}$ ,  $B[h_i(x)] = 1$
- *Lookup* ( $B, x$ )
  - If  $\bigwedge_{i \in \{1, \dots, k\}} B[h_i(x)]$  , return PRESENT, else ABSENT




# Bloom Filter

- If the element  $x$  has been added to the Bloom filter, then  $Lookup(B, x)$  always return PRESENT



# Bloom Filter

- If the element  $x$  has been added to the Bloom filter, then  $Lookup(B, x)$  always return PRESENT
- If  $x$  has not been added to the filter before?
  - $Lookup$  sometimes still return PRESENT



# Designing Bloom Filter

- Want to minimize the probability that we return a false positive
- Parameters  $m = |B|$  and  $k =$  number of hash functions
- $k = 1 \Rightarrow$  normal bit-array
- What is effect of changing  $k$ ?



# Effect of number of hash functions

- Increasing  $k$ 
  - Possibly makes it harder for false positives to happen in *Lookup* because of  $\bigwedge_{i \in \{1, \dots, k\}} B[h_i(x)]$
  - But also increases the number of filled up positions
- We can analyse to find out an “optimal  $k$ ”

# False positive analysis

- $m = |B|$ ,  $n$  elements inserted
- If  $x$  has not been inserted, what is the probability that  $Lookup(B, x)$  returns PRESENT?

# False positive analysis

- $m = |B|$ ,  $n$  elements inserted
- If  $x$  has not been inserted, what is the probability that  $Lookup(B, x)$  returns PRESENT?
- Assume  $\{h_1, h_2, \dots, h_k\}$  are independent and  $\Pr[h_i(\cdot) = j] = \frac{1}{m}$  for all positions  $j$
- $\Pr[h_i(x) = 0] = \left(1 - \frac{1}{m}\right)^{kn} \approx e^{-kn/m}$

# False positive analysis

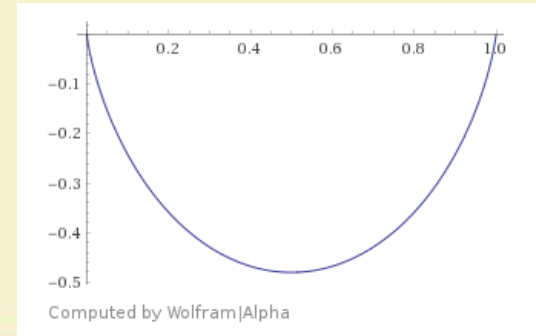
- The expected number of zero bits  $\approx me^{-kn/m}$  w.h.p.
- $\Pr[\text{Lookup}(B, x) = \text{PRESENT}] = (1 - e^{-kn/m})^k$
- Can we choose  $k$  to minimize this probability

# Choosing number of hash functions

- $p = e^{-kn/m}$
- Log (False Positive) =

$$\log(1 - p)^k = k \log(1 - p) = -\frac{m}{n} \log(p) \log(1 - p)$$

Minimized at  $p = \frac{1}{2}$ , i.e.  $k = m \log(2)/n$





# Bloom filter design

- This “optimal” choice gives false positive =  $2^{-m \log(2)/n}$
- If we want a false positive rate of  $\delta$  , set  $m = \left\lceil \frac{\log\left(\frac{1}{\delta}\right)n}{\log^2(2)} \right\rceil$

Example: If we want 1% FPR, we need 7 hash functions and total  $10n$  bits

# Applications

- Widespread applications whenever small false positives are tolerable
- Used by browsers
  - to decide whether an URL is potentially malicious: a BF is used in browser, and positives are actually checked with the server.
- Databases e.g. BigTable, HBase, Cassandra, Postgrepsql use BF to avoid disk lookups for non-existent rows/columns
- Bitcoin for wallet synchronization....

# Handling deletions

- Chief drawback is that BF does not allow deletions
- Counting Bloom Filter [Fan et al 00]
  - Every entry in BF is a small counter rather than a single bit
  - $Insert(x)$  increments all counters for  $\{h_i(x)\}$  by 1
  - $Delete(x)$  decrements all  $\{h_i(x)\}$  by 1
  - maintains 4 bits per counter
  - False negatives can happen, but only with low probability



# Other Extensions

- Many recent work on Bloom filters
  - Can we do with less hashing?
  - Can BFs be compressed (needed for distributed systems)
  - Are there better structures that use less space, less randomness and less memory lookups?



# References:

- Primary reference for this lecture
  - Survey on Bloom Filter, Broder and Mitzenmacher 2005,  
<https://www.eecs.harvard.edu/~michaelm/postscripts/im2005b.pdf>
  - <http://www.firatatagun.com/blog/2016/09/25/bloom-filters-explanation-use-cases-and-examples/>
- Others
  - Randomized Algorithms by Mitzenmacher and Upfal.

# Thank You!!



**IIT Gandhinagar**  
Indian Institute of  
Technology Gandhinagar



NPTEL ONLINE  
CERTIFICATION COURSES

Anirban Dasgupta  
Computer Science and Engg.