



IIT KHARAGPUR
IIT GANDHINAGAR



NPTEL ONLINE
CERTIFICATION COURSES

Scalable Data Science

Lecture 13b: Data Dependent LSH


Anirban Dasgupta

Computer Science and Engineering
IIT GANDHINAGAR



IIT Gandhinagar
Indian Institute of
Technology Gandhinagar

Locality Sensitive Hashing



Given input data, radius r , approx factor c and confident δ

Output: if there is any point at distance $\leq r$ then w.p. $1 - \delta$ return one at distance $\leq cr$

Algo: Choose (k, L) .

do L times

iid hash functions : $\{h_{i1} \dots h_{ik}\}$

Create hash table H_i by putting each x in bucket

$$H_i(x) = (h_{i1}(x), \dots, h_{ik}(x))$$

Store non-empty buckets in normal hash table

Picture courtesy Slaney et al.

Issues

- Parameters k , L need to be tuned for each domain
- Random directions are meant to create a random partitioning of the dataset
- While useful to guard against “worst case datasets”, we do not exploit the dataset structure

Hashing as binary codes

- Assume points are in Euclidean space
- How can we get binary vectors so that Hamming distance approximates Euclidean distance

Properties of a binary code

- Should be easily computable
- Should preserve distances approximately
- Should have small number of bits
 - the bits should be independent and unbiased

Optimization

- W_{ij} = similarity between i and j
 - Say $W_{ij} = \exp\left(-\frac{|x_i - x_j|^2}{s}\right)$
- y_i = codeword for point i
- $|y_i - y_j|^2$ also equals Hamming(i, j)

Learning codes

- Average hamming distance = $\sum_{ij} W_{ij} |y_i - y_j|^2$
- $y_i \in \{-1, +1\}^k$
- Each bit should be unbiased: $\sum_i y_i = 0$
- Bits should be uncorrelated $\sum_i y_i y_i^t = I$



Casting as optimization problem

[Waiss et al.]

- Can we solve : minimize $\sum_{ij} W_{ij} |y_i - y_j|^2$
- subject to
 - $y_i \in \{-1, +1\}^k$
 - $\sum_i y_i = 0$
 - $\sum_i y_i y_i^t = I$



Hardness

- Unfortunately, no!, even for single bit
- Graph partitioning problem: For graph G partition $V(G)$ into two sets A and B such that $|A| = |B|$ and

$$\text{minimize } \sum_{i \in A, j \in B} W_{ij}$$



Spectral Relaxation

- $Y = n \times k$ code matrix
- Diagonal D , $D_{ii} = \sum_j W_{ij}$
- minimize $\sum_{ij} W_{ij} |y_i - y_j|^2 = \text{trace}(Y^t (D - W) Y)$
 - $Y^t \cdot \mathbf{1} = 0$
 - $Y^t Y = I$
 - Drop the constraint that Y are in $\{-1, +1\}$

Spectral codes

- The above problem is solved by $Y =$ smallest k eigenvectors of $D - W$
 - After dropping the one with value 0
- To get codes,
 - We could threshold eigenvectors, but then hard to extend it for query

Eigenvectors

- Assume that the data is coming from some distribution in R^d
 - But estimating this distribution is hard also
 - We could try to interpolate the eigenvectors to query points, under above assumptions, but is computationally expensive (Nystrom extension)

Eigenvectors

- Assume that the data is coming from some distribution in R^d
 - But estimating this distribution is hard also
 - We could try to interpolate the eigenvectors to query points, under above assumptions, but is computationally expensive (Nystrom extension)
- Assume data distribution is product of uniform distributions
 - Use PCA to find the axes




Eigenfunctions

- Take limit of eigenvectors as $n \rightarrow \infty$, and consider the “normalized” similarity matrix (Laplacian)
- Analytical form of Eigenfunctions exists for certain distributions (uniform, Gaussian)

- For uniform

$$\begin{aligned}\Phi_k(x) &= \sin\left(\frac{\pi}{2} + \frac{k\pi}{b-a}x\right) \\ \lambda_k &= 1 - e^{-\frac{\epsilon^2}{2} \left|\frac{k\pi}{b-a}\right|^2}\end{aligned}$$




[Image from Weiss et al]

- Constant time calculation for any new point


Algorithm

Input: Data $\{x_i\}$, target dimensionality k



Algorithm

Create top k PCA of $D - W$



Gives us top k axes

Find the $[a_i, b_i]$ for each axes

and create $\phi_1(x) \dots \phi_k(x)$ for each direction

Algorithm

Create top k PCA of $D - W$


Gives us top k axes

Find the $[a_i, b_i]$ for each axes

and create $\phi_{i1}(x) \dots \phi_{ik}(x)$


and $\lambda_{i1} \dots \lambda_{ik}$ for each direction

Total dk eigenvalues \rightarrow sort and take the top k eigenvalues and corresponding functions




Algorithm

Threshold chosen
Eigenfunctions




Empirical observation: bit codes
seem robust to the uniform assumption



Results

- Shown to have better properties than naïve LSH on large datasets



[Image from Weiss et al]

Summary

- Large literature on learning the hash codes rather than use random projection
 - many many ways to learn from data
- Unfortunately, theoretical guarantees are not available for such data-dependent version
 - time to calculate projections might also be higher
- Recent papers point to analysis techniques that can bridge theory and practice



References:

- Primary references for this lecture
 - Spectral Hashing, Yair Weiss, Antonio Torralba and Rob Fergus. [*NIPS*], 2008

Thank You!!



IIT Gandhinagar
Indian Institute of
Technology Gandhinagar



NPTEL ONLINE
CERTIFICATION COURSES

Anirban Dasgupta
Computer Science and Engg.