



IIT KHARAGPUR
IIT GANDHINAGAR



NPTEL ONLINE
CERTIFICATION COURSES

Scalable Data Science

Lecture 14c: Fast LSH + Sparse Random Projection

Anirban Dasgupta

Computer Science and Engineering

IIT GANDHINAGAR






IIT Gandhinagar

Indian Institute of
Technology Gandhinagar

Johnson Lindenstrauss Lemma [JL84]

$\epsilon > 0$, $k \geq \frac{C}{\epsilon^2} \log(n)$. There exists a **linear mapping** A such that whp, for all (i, j)

$$(1 - \epsilon)|x_i - x_j| \leq |Ax_i - Ax_j| \leq (1 + \epsilon)|x_i - x_j|$$



$$x_i \in \mathbb{R}^d$$

Time taken for projection

- Projection matrix out of Gaussian or iid $\pm 1 = O(kd)$

- $k = \Omega\left(\frac{1}{\epsilon^2}\right)$

- FJLT : projection matrix is $\frac{1}{\sqrt{d}} PHD$

- H is the $d \times d$ Hadamard matrix
 - D is a random ± 1 diagonal matrix
 - P is a sparse Gaussian matrix
 - Time = $O(d \log d + k \log(nd))$

Notice the
normalization $\frac{1}{\sqrt{d}}$



Densification claim

$$x \in \mathbb{R}^d, \|x\|_2 = 1$$

Claim: $\max_i |(H D x)_i| \leq O\left(\frac{\log(nd)}{d}\right)^{1/2}$

Application of Chernoff style tail inequality per coordinate and union bound

Projecting a dense vector

$$y = HDx, \max_i |y_i| \approx O(\sqrt{\log(nd)/d})$$

$$P = \begin{cases} 0, w.p. 1 - q \\ N\left(0, \frac{1}{\sqrt{q}}\right), w.p. q \end{cases}, P \in \mathbb{R}^{k \times d}, k = O\left(\frac{1}{\epsilon^2} \log\left(\frac{1}{\delta}\right)\right)$$


$z = \frac{1}{\sqrt{d}} PHDx$ is the final projected vector

Fast JL Transform [AC09]

If $q = O(|x|_\infty^2) = O\left(\frac{\log(nd)}{d}\right)$, $\frac{1}{\sqrt{d}}$ PHD satisfies JL property

Calculating $y = PHDx$ takes time $O(d \log d + k \log(nd))$, potentially much faster than original Gaussian construction

Locality Sensitive Hashing



Given input data, radius r , approx factor c and confident δ

Output: if there is any point at distance $\leq r$ then w.p. $1 - \delta$ return one at distance $\leq cr$

$$h_i(x) = \left\lfloor \frac{x \cdot v_i + b_i}{w} \right\rfloor$$

Picture courtesy Slaney et al.

Time taken

- Total query time = time to hash + time to check all candidates
 - Calculating k-hash indices takes time $O(kd)$
 - Calculating indices for L buckets takes time $O(kdL)$
- Can we reduce query time?



Creating hash indices

- Looking at LSH as random projection + quantization
 - $A \in R^{k \times d}$ is a Gaussian JL matrix, $b \in R^k$
 - We first project and then bucketize
 - calculate $\left\lfloor \frac{Ax+b}{w} \right\rfloor$, k-index key calculated at once
- Time taken by matrix-vector multiplication = $O(kd)$ per hash table

Collision Probability

- $p(u) = \Pr[h_i(p) = h_i(q)]$ when $|p - q| = u$
- $p(u) = \int_0^w \frac{1}{u} f\left(\frac{t}{u}\right) \left(1 - \frac{t}{w}\right) dt$, $f(v) = \text{pdf of } |N(0,1)|$
- This is decreasing with increasing u



ACHash_[DKS11]

- When calculating a k-tuple hash bucket index
 - $\left\lfloor \frac{Ax+b}{w} \right\rfloor$, use $A = PHD$
 - $q \approx O\left(\frac{\log(d)}{d}\right)$
 - Projection time = $O(d \log d + kL \log^2 d)$

ACHash

$p_{AC}(u)$ = probability that a k -tuple hash bucket has same value for two points at distance u

We can show that

$$-(k+1)\delta + p^k((1+\epsilon)u) \leq p_{AC}(u) \leq p^k((1-\epsilon)u) + (k+1)\delta.$$

i.e. collision prob does not change much

DHHash

We can make the projection faster

D = random diagonal matrix of ± 1

G = random diagonal matrix, $G_{ii} \sim N(0, 1)$

M = random permutation matrix


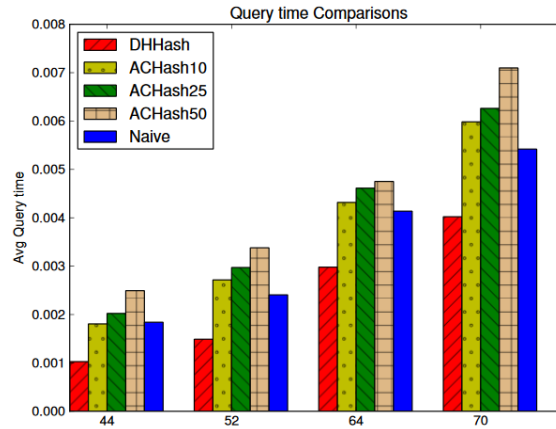
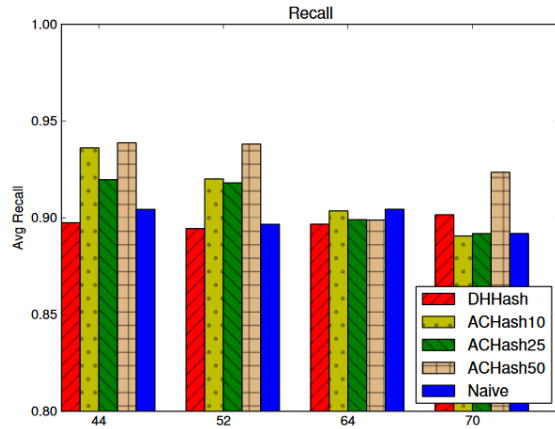
Hash value calculated as $\left\lfloor \frac{HGMHDx + b}{w} \right\rfloor$

DHHash

- The above creates d bits
- Sample kL indices and create L hash bucket ids, each of size k
- Total calculation time for all L bucket-ids = $O(d \log d + kL)$
- Also performs nicely in practice



Experiments: faster query time with more or less same recall



(d) Recall, LSH query time, and LSH computation time for P53.

FJLT is fast, but...

- Sparse vectors are prevalent in large scale ML
 - e.g. document representations
- What happens when a sparse vector is
 - multiplied by a dense Gaussian matrix
 - multiplied by *PHD* of FJLT



Effect on sparsity

- Sparse vectors are prevalent in large scale ML
 - e.g. document representations
- What happens when a sparse vector is
 - multiplied by a dense Gaussian matrix
 - multiplied by *PHD* of FJLT
- Both result in dense vector!
 - much more expensive in terms of storage and computation



Preserving sparsity

- Can we design a linear transformation $A \in \mathbb{R}^{k \times d}$ such that
 - $|Ax| \approx |x|$ w.h.p. for any fixed $x \in \mathbb{R}^d$
 - $\text{nnz}(x) \approx \text{nnz}(Ax)$
- Existing iid constructions do not satisfy this
 - As we saw before, we cannot make the projection matrix very sparse if the elements are chosen independently

Hashing as projection


Input $x \in \mathbb{R}^d$

Target $y \in \mathbb{R}^k$

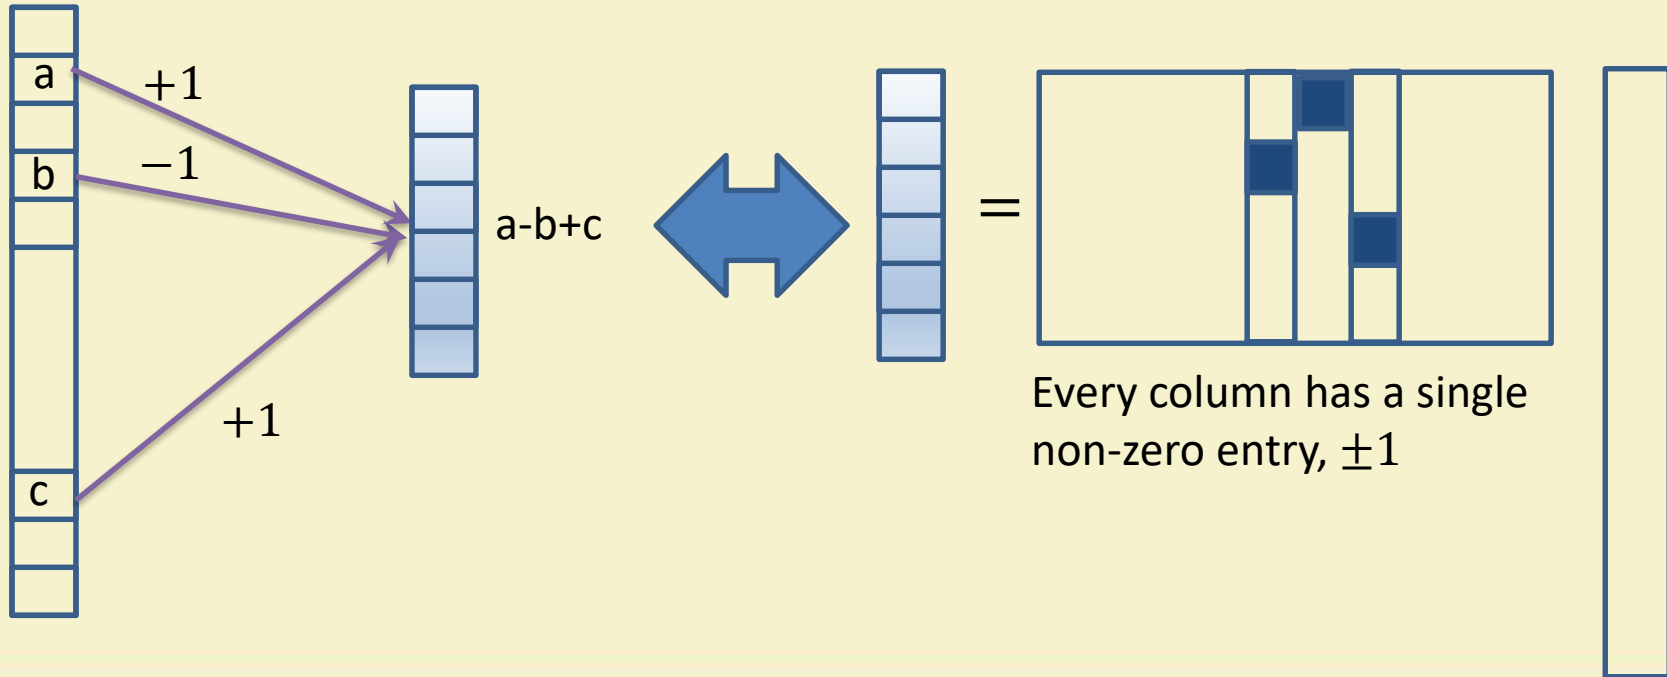
Hash function $h: [d] \rightarrow [k]$

Sign hash: $s: [h] \rightarrow \{-1, +1\}$

$$y[j] = \sum_{i:h(i)=j} s(i) x_i$$



Hashing as projection



Sparsity

- $nnz(y) = nnz(x)$
- Norm preservation whp does not happen with only one hash function
 - Repeat the construction



Sparse random projection matrix



For every column, choose a fixed number, ℓ , positions
For each position chosen, fill up with i.i.d. ± 1 random variable

Formalization

$$A \in \mathbb{R}^{k \times d}$$

[DKS10] Choose ℓ positions from each column with replacement

[KN11] Choose ℓ positions without replacement

For each nonzero position $A_{ij} = \begin{cases} +1 & \text{w.p. } \frac{1}{2} \\ -1 & \text{w.p. } \frac{1}{2} \end{cases}$

Guarantee

Claim: For $\ell = \tilde{O}\left(\frac{1}{\epsilon}\right)$, $k = O\left(\frac{1}{\epsilon^2} \log\left(\frac{1}{\delta}\right)\right)$,
 $\Pr[(1 - \epsilon) \leq |Ax| \leq (1 + \epsilon)] \geq 1 - \delta$

So a vector that initially has $nnz(x)$ nonzeros, now will have at most $\frac{nnz(x)}{\epsilon}$ non-zeros



References:

- Primary references for this lecture
 - Fast Johnson Lindenstrauss Transform, Ailon and Chazelle, SIAM J Computing 2009.
 - Sparse Johnson Lindenstrauss Transformation, Dasgupta, Kumar, Sarlos, STOC 2010.
 - Fast Locality Sensitive Hashing, Dasgupta, Kumar, Sarlos, KDD 2011.

Summary

In this lecture, we saw

- An application of FJLT in improving the query time of LSH
- A different construction for random projection that, in addition to preserving length, preserves sparsity better than original

Thank You!!



IIT Gandhinagar
Indian Institute of
Technology Gandhinagar



NPTEL ONLINE
CERTIFICATION COURSES

Anirban Dasgupta
Computer Science and Engg.