



IIT KHARAGPUR  
IIT GANDHINAGAR



NPTEL ONLINE  
CERTIFICATION COURSES

# Scalable Data Science

## Lecture 18: Spark

Sourangshu Bhattacharya  
Computer Science and Engineering  
IIT KHARAGPUR

# In the previous lectures:

- Outline:
  - Scala
    - Var and Val
    - Classes and objects
    - Functions and higher order functions
    - Lists

# In this Lecture:

- Outline:
  - Spark
    - Motivation
    - RDD
    - Actions and transformations
    - Examples:
      - Matrix multiplication
      - Logistic regression
      - Pagerank

# SPARK



IIT KHARAGPUR



NPTEL ONLINE  
CERTIFICATION COURSES

# Spark

**Spark is an In-Memory Cluster Computing platform for Iterative and Interactive Applications.**

<http://spark.apache.org>

# Spark

- ☐ Started in AMPLab at UC Berkeley.
- ☐ Resilient Distributed Datasets.
- ☐ Data and/or Computation Intensive.
- ☐ Scalable – fault tolerant.
- ☐ Integrated with SCALA.
- ☐ Straggler handling.
- ☐ Data locality.
- ☐ Easy to use.

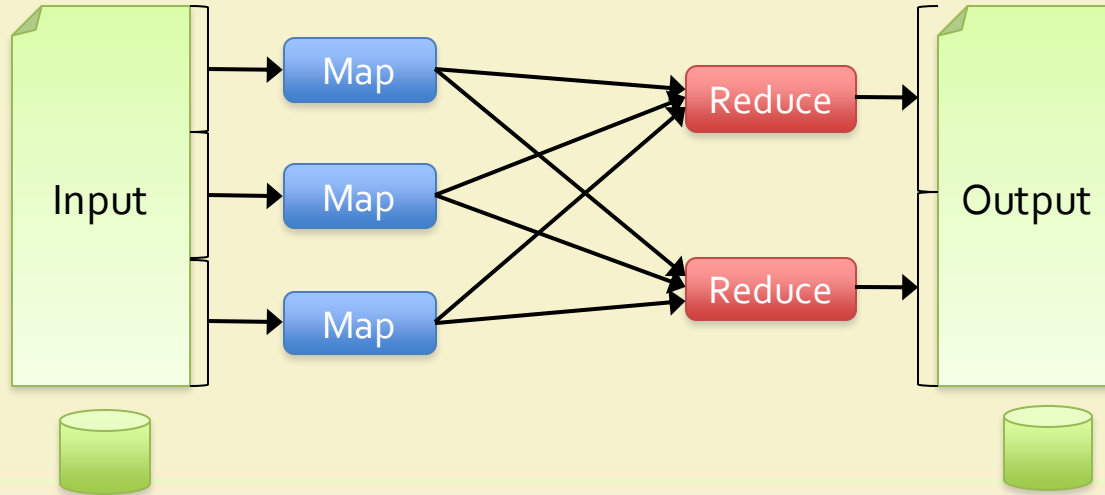
# Background

- Commodity clusters have become an important computing platform for a variety of applications
  - **In industry:** search, machine translation, ad targeting, ...
  - **In research:** bioinformatics, NLP, climate simulation, ...
- High-level cluster programming models like MapReduce power many of these apps
- *Theme of this work: provide similarly powerful abstractions for a broader class of applications*

# Motivation

Current popular programming models for clusters transform data flowing from stable storage to stable storage

E.g., MapReduce:





# Motivation

- Current popular programming models for clusters transform data flowing from stable storage to stable storage
- E.g., MapReduce:

**Benefits of data flow:** runtime can decide where to run tasks and can automatically recover from failures

# Motivation

- Acyclic data flow is a powerful abstraction, but is not efficient for applications that repeatedly reuse a *working set* of data:
  - **Iterative** algorithms (many in machine learning)
  - **Interactive** data mining tools (R, Excel, Python)
- Spark makes working sets a first-class concept to efficiently support these apps

# Spark Goal

- Provide distributed memory abstractions for clusters to support apps with working sets
- Retain the attractive properties of MapReduce:
  - Fault tolerance (for crashes & stragglers)
  - Data locality
  - Scalability

**Solution:** augment data flow model with “resilient distributed datasets” (RDDs)

# Resilient Distributed Datasets

- ☐ Immutable distributed SCALA collections.
  - ☐ Array, List, Map, Set, etc.
- ☐ Transformations on RDDs create new RDDs.
  - ☐ Map, ReduceByKey, Filter, Join, etc.
- ☐ Actions on RDD return values.
  - ☐ Reduce, collect, count, take, etc.
- ☐ Seamlessly integrated into a SCALA program.
- ☐ RDDs are materialized when needed.
- ☐ RDDs are cached to disk – graceful degradation.
- ☐ Spark framework re-computes lost splits of RDDs.

# RDDs in More Detail

- ❑ An RDD is an immutable, partitioned, logical collection of records
  - ❑ Need not be materialized, but rather contains information to rebuild a dataset from stable storage
- ❑ Partitioning can be based on a key in each record (using hash or range partitioning)
- ❑ Built using bulk transformations on other RDDs
- ❑ Can be cached for future reuse

# RDD Operations

## Transformations (define a new RDD)

map  
filter  
sample  
union  
groupByKey  
reduceByKey  
join  
cache

...

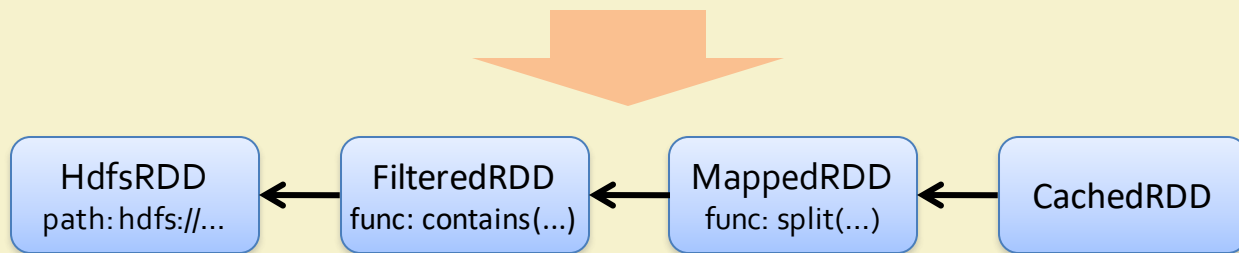
## Actions (return a result to driver)

reduce  
collect  
count  
save  
lookupKey  
...

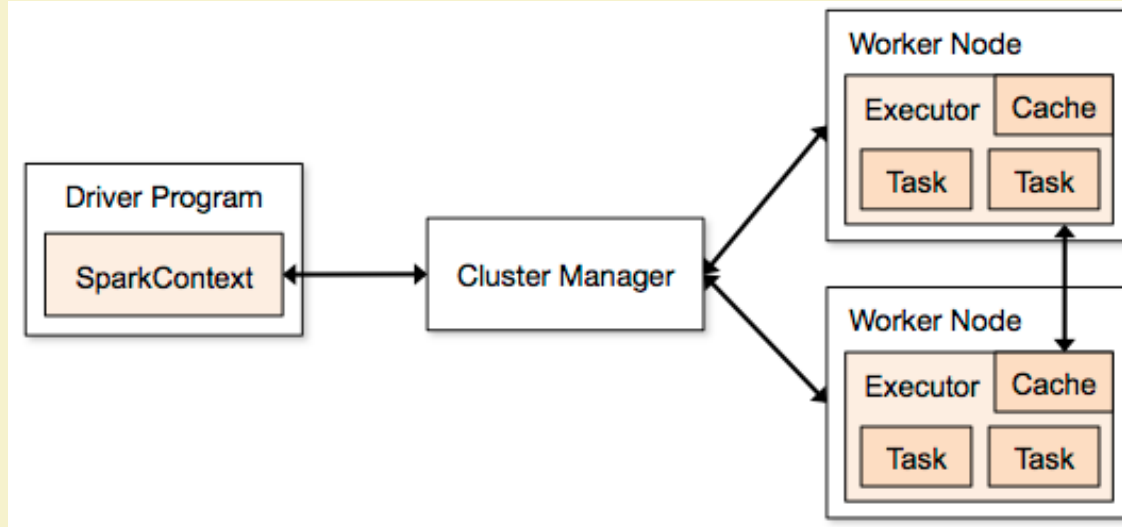
# RDD Fault Tolerance

- RDDs maintain *lineage* information that can be used to reconstruct lost partitions
- Ex: 

```
cachedMsgs = textFile(...).filter(_.contains("error"))  
                                .map(_.split('\t')(2))  
                                .cache()
```



# Spark Architecture





# Example: MapReduce

- MapReduce data flow can be expressed using RDD transformations

```
res = data.flatMap(rec => myMapFunc(rec))  
            .groupByKey()  
            .map((key, vals) => myReduceFunc(key, vals))
```

Or with combiners:

```
res = data.flatMap(rec => myMapFunc(rec))  
            .reduceByKey(myCombiner)  
            .map((key, val) => myReduceFunc(key, val))
```

# Word Count in Spark

```
val lines = spark.textFile("hdfs://...")  
  
val counts = lines.flatMap(_.split("\\s"))  
                    .reduceByKey(_ + _)  
  
counts.save("hdfs://...")
```

# Example: Matrix Multiplication

# Matrix Multiplication

- ◆ Representation of Matrix:
  - ◆ List <Row index, Col index, Value>
  - ◆ Size of matrices: First matrix (A):  $m \times k$ , Second matrix (B):  $k \times n$
- ◆ Scheme:
  - ◆ For each input record: If input record
- ◆ Mapper key: <row\_index\_matrix\_1, Column\_index\_matrix\_2>
- ◆ Mapper value: < column\_index\_1 / row\_index\_2, value>
- ◆ GroupByKey: List(Mapper Values)
- ◆ Collect all (two) records with the same first field multiply them and add to the sum.

# Example: Logistic Regression



IIT KHARAGPUR



NPTEL ONLINE  
CERTIFICATION COURSES

# Logistic Regression

- Binary Classification.  $y \in \{+1, -1\}$
- Probability of classes given by linear model:

$$p(y | x, w) = \frac{1}{1 + e^{(-yw^T x)}}$$

- Regularized ML estimate of  $w$  given dataset  $(x_i, y_i)$  is obtained by minimizing:

$$l(w) = \sum_i \log(1 + \exp(-y_i w^T x_i)) + \frac{\lambda}{2} w^T w$$

# Logistic Regression

- Gradient of the objective is given by:

$$\nabla l(w) = \sum_i (1 - \mathcal{S}(y_i w^T x_i)) y_i x_i - / w$$

- Gradient Descent updates are:

$$w^{t+1} = w^t - s \nabla l(w^t)$$

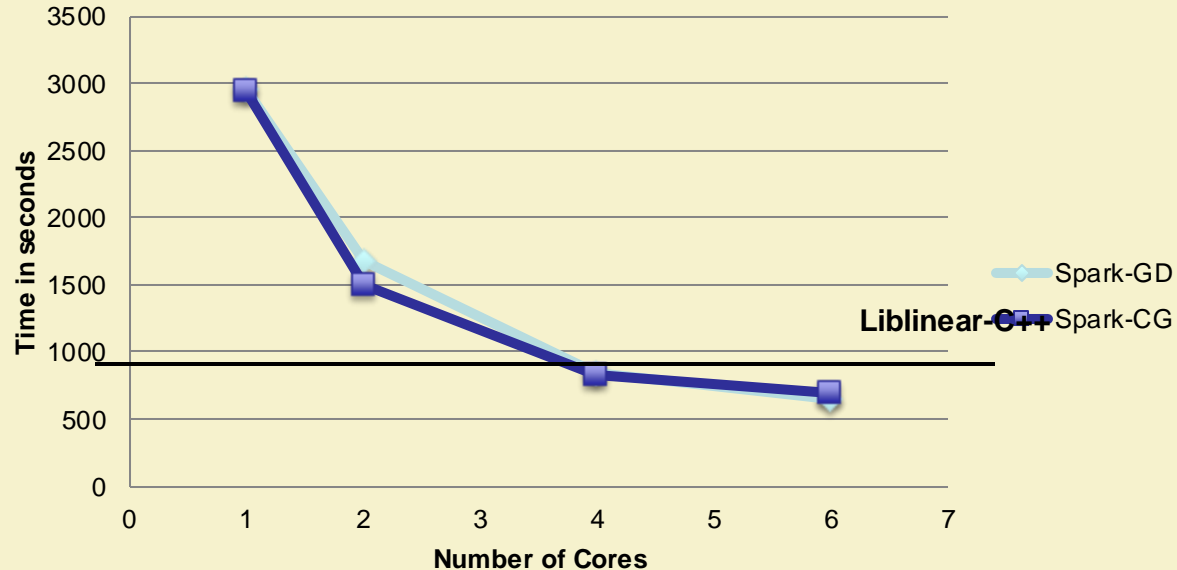
# Spark Implementation

```
val x = loadData(file) //creates RDD
var w = 0
do {
  //creates RDD
  val g = x.map(a => grad(w,a)).reduce(_+_ )
  s = linesearch(x,w,g)
  w = w - s * g
}while(norm(g) > e)
```



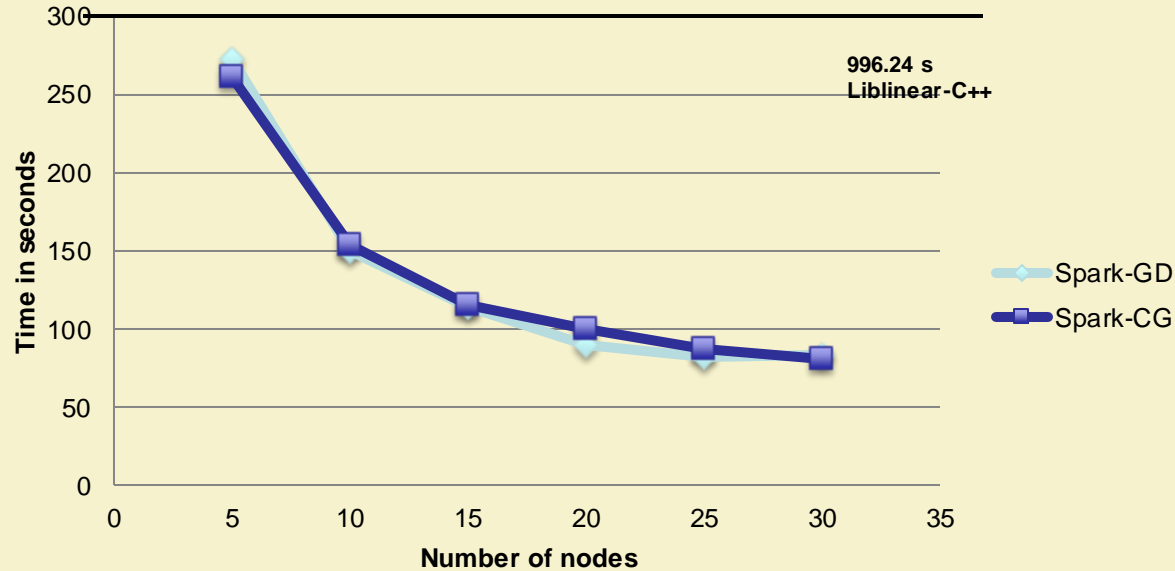
# Scaleup with Cores

## Epsilon (Pascal Challenge)



# Scaleup with Nodes

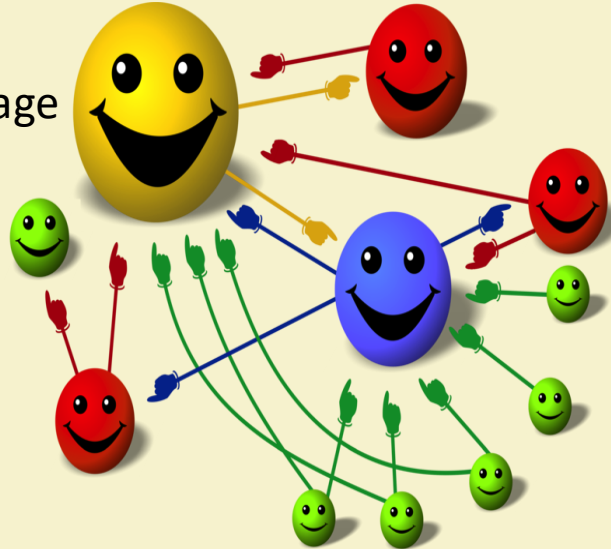
## Epsilon (Pascal Challenge)



# Example: PageRank

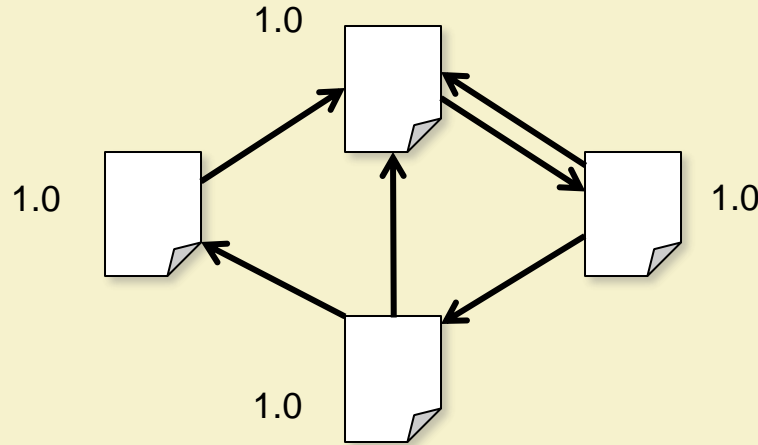
# Basic Idea

- Give pages ranks (scores) based on links to them
  - Links from many pages  
→ high rank
  - Link from a high-rank page  
→ high rank



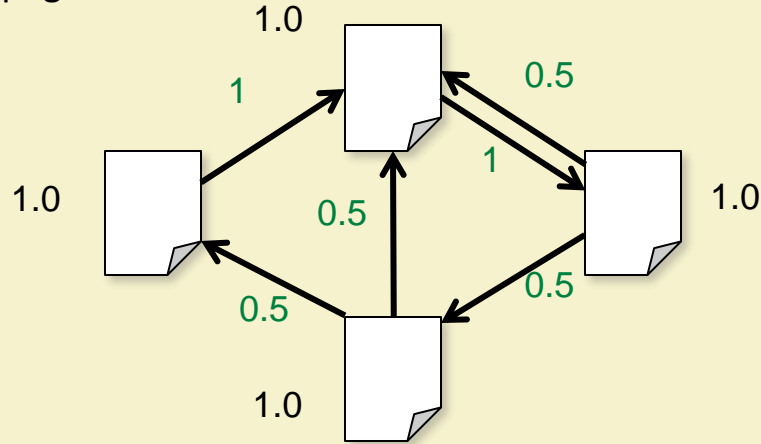
# Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page  $p$  contribute  $\text{rank}_p / |\text{neighbors}_p|$  to its neighbors
3. Set each page's rank to  $0.15 + 0.85 \times \text{contribs}$



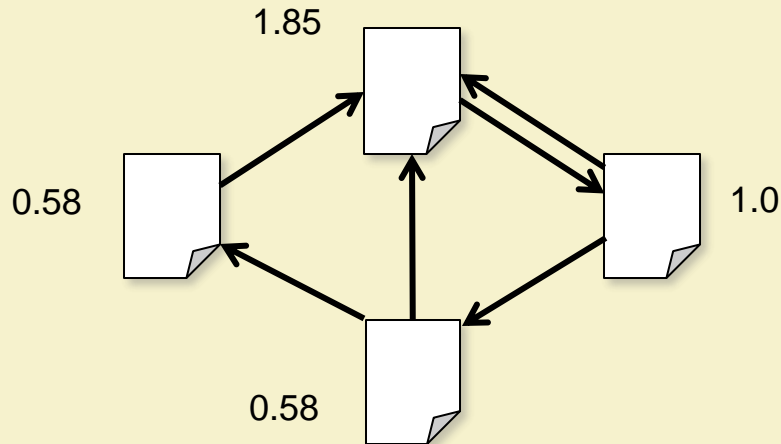
# Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page  $p$  contribute  $\text{rank}_p / |\text{neighbors}_p|$  to its neighbors
3. Set each page's rank to  $0.15 + 0.85 \times \text{contribs}$



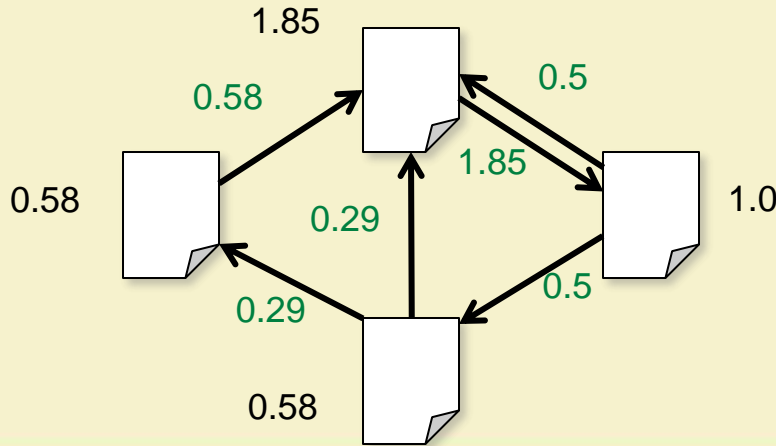
# Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page  $p$  contribute  $\text{rank}_p / |\text{neighbors}_p|$  to its neighbors
3. Set each page's rank to  $0.15 + 0.85 \times \text{contribs}$



# Algorithm

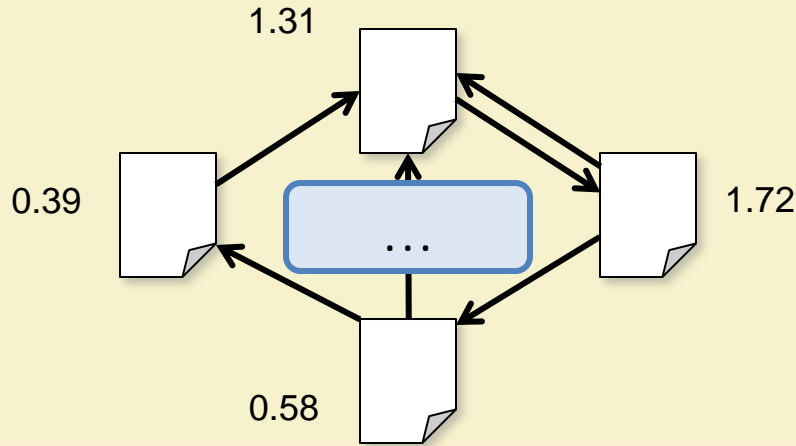
1. Start each page at a rank of 1
2. On each iteration, have page  $p$  contribute  $\text{rank}_p / |\text{neighbors}_p|$  to its neighbors
3. Set each page's rank to  $0.15 + 0.85 \times \text{contribs}$





# Algorithm

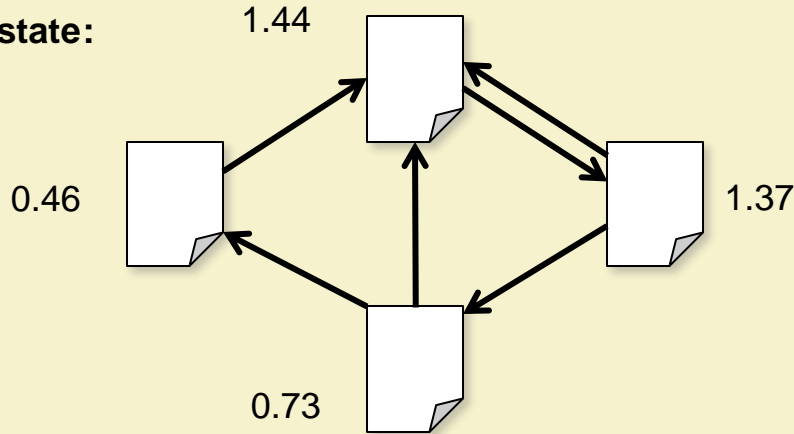
1. Start each page at a rank of 1
2. On each iteration, have page  $p$  contribute  $\text{rank}_p / |\text{neighbors}_p|$  to its neighbors
3. Set each page's rank to  $0.15 + 0.85 \times \text{contribs}$



# Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page  $p$  contribute  $\text{rank}_p / |\text{neighbors}_p|$  to its neighbors
3. Set each page's rank to  $0.15 + 0.85 \times \text{contribs}$

Final state:



# Spark Implementation

```
val links = // RDD of (url, neighbors) pairs
var ranks = // RDD of (url, rank) pairs

for (i <- 1 to ITERATIONS) {
  val contribs = links.join(ranks).flatMap {
    (url, (nhb, rank)) =>
      nhb(dest => (dest, rank/nhb.size))
  }
  ranks = contribs.reduceByKey(_ + _)
                    .mapValues(0.15 + 0.85 * _)
}

ranks.saveAsTextFile(...)
```

# Conclusion:

- We have seen:
  - Spark
    - Motivation
    - RDD
    - Actions and transformations
    - Examples:
      - Matrix multiplication
      - Logistic regression
      - Pagerank

# References:

- Learning Spark: Lightning-Fast Big Data Analysis. Holden Karau, Andy Konwinski, Patrick Wendell, Matei Zaharia. O Reilly Press 2015.
- Any book on scala and spark.

# Thank You!!



IIT KHARAGPUR



NPTEL ONLINE  
CERTIFICATION COURSES

Sourangshu Bhattacharya  
Computer Science and Engg.